

Fundamentals of Artificial Intelligence (57205HT16): Assignment 2 - Happy, Sad, Mischievous or Mad?

Marc Coquand	Linus Lagerhjelm
id14mcd	id14llm
mcoquand@gmail.com	id14llm@cs.umu.se

Supervisor: Thomas Johansson
Alexander Sutherland
Thomas Hellström

October 10, 2016

Contents

1	Introduction	3
1.1	Running the program	3
2	Problem description	3
2.1	Perceptron Learning Algorithm	4
3	Concluding discussion	6

1 Introduction

This report contains the documentation for a perceptron based classification system that guesses the emotional state of faces presented as input. The inputs are 20 x 20 pixel maps with 32 grey levels from white to black.

1.1 Running the program

The implementation is written in python 2.7 which means that a python 2.7 shell is required to run the implementation. Required program arguments are a search path to a `training.txt`, search path to an answer file `training-facit.txt` and a test file `test-file.txt`. The format of the training file should be the id of the image followed by the image on a new line. The answer file should be formatted as the image id follow by the correct answer (1,2,3 or 4).

Should the user fail to provide the required arguments, the program will immediately terminate with a message explaining that not enough input arguments were provided.

Example usage of the program might look like:

```
python faces.py ../data/FaceTest/training-file.txt
../data/FaceTest/training-facit.txt test-file.txt
```

2 Problem description

The program works as such that it starts of by reading the *images*. An image consists of it's pixels (value from 0 - white to 31 - black) and it's ID (sad, mischievous, happy, mad). The program then filters the noise and then sends the images in the *Tutor* class. The tutor uses Perceptrons to learn what images look like (algorithm explained below). A perceptron is an algorithm for learning a binary classifier. Afterwards the Examiner class makes a guess and predicts which image belongs to which on a new set of image. The program will then output the percentage of the guesses that were correct.

2.1 Perceptron Learning Algorithm

The program implements an algorithm for supervised learning called *Perceptron Learning*. *Supervised Learning* is a technique for machine learning where the Artificial neural network is exposed to input and produces a guess based on the knowledge acquired from previous guesses. Since the desired output from the ANN, in supervised learning, is known for every input, the ANN is slightly modified with each iteration to eventually produce good outputs for every set of inputs.

The ANN in the produced program is composed by four different Perceptrons where each one represents one of the different facial expressions that are to be identified in the assignment.

Each perceptron is implemented in *Python* as a class composed by a list of 401 weights that maps to the 400 pixel values that makes up an image. The additional weight is provided as a bias for the perceptron. A bias is used in order to shift the result of the activation function (more on that later) towards either the left or the right. I.e. to make the output more extreme so that it will produce a more distinct answer.

Training in the network occurs in the *Tutor* class where the provided set of training images is repeatedly shown for each perceptron. The output from a perceptron, also known as *activation*, is computed using the formula:

$$a_i = act(\sum_j x_j * w_{j,i})$$

where x is the input, w is the weight and act is the activation function, which is the sigmoid function:

$$sigmoid(x) = 1/(1 + e^{-x})$$

The difference between the expected output and the actual output is the error e . The error e is then multiplied with the learning rate α and the pixel matrix to produce the gradient ∇w . This gradient is then added to every weight in the

internal weight matrix for each perceptron. This course of events can be described using the bellow formulas:

$$\begin{aligned}e_i &= y_i - a_i \\ \nabla w_{j,i} &= \alpha e_i x_j \\ w_{j,i} &\leftarrow w_{j,i} + \nabla w_{j,i}\end{aligned}$$

This process is repeated for the entire set of training images, however, only performing this process once would not produce very good results and therefore, when the network has used all the training data, the images are shuffled and the entire process described above is repeated.

The main problem encountered while training an ANN is finding the fine line right before the network gets overfitted. When the ANN gets overfitted, it becomes too specialized on the training data which will render it useless when exposed to previously unseen data. It is, however, still crucial to train the ANN as close to this limit as possible in order for it to be able to make generalizations to the data and produce decent output.

One way of determine when the network has reached the desired level of training is, for each iteration of the training data set, use a formula to compute the sum of the output errors for the network and end training when the value from this function gets below a predefined level. The function used in the current implementation is:

$$E = \frac{1}{2} \sum_{i=1}^p ||y^{(i)} - d^{(i)}||^2$$

Where p is the number of input/output vector pairs in the training set.

When the threshold for the error is reached, the perceptrons is deemed ready to be passed on to the *Examiner* class where the network is exposed to the set of test images. Each image is shown to every perceptron in the network and the activation for each perceptron is recorded. The largest activation is chosen as the answer for

the network and is recorded into our list of answers. When every image in the test set has been classified the answers are printed to *stdout*.

3 Concluding discussion

The biggest difficulty was choosing which programming language to use. It was difficult choosing between Haskell and Python. Ultimately we chose Python as our knowledge in Haskell was not sufficient to write this program. It would be interesting to implement more AI and more specifically in machine learning programs in a functional programming language as it's easier to deal with concurrency. If you wanted to scale this program and have a larger neural network it might be necessary to add concurrency.

```

1 import sys
import re
3 from Enum import Enum
from Perceptron import Perceptron
5 from Tutor import Tutor
from Utils import flatten
7 from Examiner import Examiner
from Image import Image
9
# Program arguments: training-set, answer-set, examination-set (
from spec)
11 EXPECTED_ARGS = 3
Types = Enum(["HAPPY", "SAD", "MISCHIEVOUS", "MAD"])
13
# Optimal values (derived from experiments):
# learning_rate = 0.01
# threshold = 1
15 LEARNING_RATE = 0.1
17 THRESHOLD = 1
19
21 def parse_ans(ans_file):
    """
23     Opens the provided answer file and extracts the answers
    which are
    returned as a list of integers. NON PURE FUNCTION!
25     :param ans_file: path to answer file
    :return: list of integers
    """
27     l2 = filter(lambda l: len(l) > 0, map(get_answer, open(
    ans_file, 'r'))))
29     return flatten(l2)
31
def parse_img_file(image_file):
    """
33     Opens the provided image file and extracts its content into
    a
35     list of image objects that gets returned. NON PURE FUNCTION!
    :return: list of image objects

```

```

37     """
38     f = open(image_file , 'r')
39     return reduce_img_file(f.readlines() , [])
40
41 def reduce_img_file(lines , images):
42     """
43     Consumes the lines of the image file in order to extract all
44     the images and their ids. In the end it will return a list
45     of image objects.
46     :return: list om image objects
47     """
48
49     if not lines: return images
50     if re.search('Image', lines[0]) is not None:
51         img = Image(extract_img_data(20, lines[1:]))
52         return reduce_img_file(lines[20:], images + [img.set_id(
53             lines[0])])
54
55     return reduce_img_file(lines[1:], images)
56
57 def extract_img_data(acc , lines):
58     """
59     Sub consumes all the lines that contains an image row and
60     returns a 2D array containing list of lists of integers
61     :param acc: int representing how many more lines to read
62     :param lines: the lines of the file
63     :return: 2D array of integers
64     """
65
66     if acc == 1: return [format_img_row(lines[0])[0]]
67     return [format_img_row(lines[acc-1])[0]] + extract_img_data(
68         acc - 1, lines)
69
70 def format_img_row(row):
71     """
72     format_img_row receives a row from the image file
73     representing a pixel
74     row in an image and returns it as a list of floats in [0,1)
75     :param row: image pixel row
76     :return: list of integers
77     """
78
79     return map(lambda l: map(lambda i: (float(i)/31), l), map(
80         lambda i: i.rstrip().split(' '), [row]))

```



```

79 def get_answer(line):
80     """
81     get_answer retrieves the value for the answer from a line in
82     the
83     answer file
84     :param line: line from answer file
85     :return: Integer representing the answer for the line
86     """
87     return map(lambda s: int(s), re.findall(r'\b\d+\b', line))

89 def validate_arguments():
90     """
91     makes sure that the user provided the right amount of
92     arguments to the
93     program, else: fail gracefully. NON PURE FUNCTION!
94     :return:
95     """
96     # Since program name is provided as additional first
97     argument
98     if len(sys.argv) < EXPECTED_ARGS + 1:
99         print "This program expects exactly %d args. %d was
100         provided" % \
101             (EXPECTED_ARGS, len(sys.argv) - 1)
102         sys.exit(0)

104 def print_results(res_arr):
105     """
106     Expects a list with formatted answer strings that gets
107     printed to stdout
108     """
109     if not res_arr: return
110     print res_arr[0]
111     print_results(res_arr[1:])

113 def main():
114     validate_arguments()
115     perceptrons = (Perceptron(Types.HAPPY), Perceptron(Types.SAD
116     ),
117                   Perceptron(Types.MISCHIEVOUS), Perceptron(
118     Types.MAD))

```

```

117     # Get content of the data files
118     img_list = parse_img_file(sys.argv[1])
119     ans_list = parse_ans(sys.argv[2])
120     test_set = parse_img_file(sys.argv[3])
121
122     # Link image and answer
123     images = map(lambda ans: img_list.pop(0).set_ans(ans),
124                  ans_list)
125
126     trained_p = Tutor(perceptrons, images, LEARNING_RATE,
127                       THRESHOLD).train()
128     print_results(Examiner(trained_p, test_set).examine())
129
130 if __name__ == "__main__":
131     main()

```

../src/faces.py

```

1 class Enum(set):
2     """ Enum is a python 2.7 implementation of a Enum
3         provided by http://stackoverflow.com/a/2182437/4689625
4     """
5     def __getattr__(self, name):
6         if name in self:
7             return name
8         raise AttributeError

```

../src/Enum.py

```

2 class Examiner:
3     """
4     Examiner is a class that, provided a set of trained neurons
5     and
6     test images, examines the training results of the network
7     Attributes:
8         perceptrons: tuple of trained perceptrons

```

```

8         test_set: set of test images
10
11     """
12     def __init__(self, perceptrons, test_set):
13         self.perceptrons = perceptrons
14         self.test_set = test_set
15
16     def examine(self):
17         """
18         Tests the provided perceptrons on the provided set of
19         images
20         :return: A list of formatted answer strings
21         """
22         return map(self._make_guess, self.test_set)
23
24     def _make_guess(self, image):
25         """
26         Exposes the provided image to the network and returns a
27         string
28         that is compliant with the specification of what a guess
29         should
30         look like
31         :return:
32         """
33         guess = self._expose_perceptrons(image, {"type": 0, "
34         activation": 0}, self.perceptrons)
35         return self._format_answer(image, guess["type"])
36
37     def _expose_perceptrons(self, image, current_guess,
38         perceptrons):
39         """
40         Shows an image to all of the perceptrons and returns the
41         largest
42         activation from the network
43         :param image: image to show
44         :param current_guess: The previous largest activation in
45         the network
46         :param perceptrons: consumed list of perceptrons
47         :return:
48         """
49         if not perceptrons: return current_guess
50         activation = perceptrons[0].process(image)
51         if activation > current_guess["activation"]:
52             current_guess = self._update_guess(activation,

```

```

46         perceptrons[0])
47         return self._expose_perceptrons(image, current_guess,
48         perceptrons[1:])
49
50     def _update_guess(self, activation, perceptron):
51         """
52         Returns a new dictionary containing the activation for
53         the
54         perceptron type
55         """
56         return {"type": perceptron.get_type_id(), "activation":
57         activation}
58
59     def _format_answer(self, img, guess):
60         """
61         Returns properly formatted string based on the networks
62         guess of
63         the provided image
64         """
65         return "" + img.get_id() + " " + str(guess)

```

../src/Examiner.py

```

1 import Utils
2 FILTER_VAL = float(3)/float(32) # Derived from experiments
3
4
5 class Image:
6     """
7     Image represents an image that holds a list of pixels, id,
8     and what
9     facial expression the image represents.
10    Attributes:
11        img: list of pixel values
12        id: the id of the image on format ImageXX
13        ans: integer 1-4 representing what facial expression the
14        img represents
15    """
16    def __init__(self, img):
17        self.img = self._preprocess_img(img)

```

```

17         self.id = None
18         self.ans = None
19
20     def _preprocess_img(self, img):
21         """
22         Performs preprocessing operations in order to sanitize
23         the input data so that it contains what we expects.
24         """
25         i1 = Utils.flatten(img)
26         i1.append(1)
27         return map(self._filter_noise, i1)
28
29     def _filter_noise(self, pix):
30         """
31         Since the images contains a lot of noise, we're
32         filtering away those pixels that won't impact the result other
33         than to make it unambiguous.
34         """
35         return pix if pix > FILTER_VAL else float(0)
36
37     def get_img(self):
38         """
39         Returns the pixel data for the image as a list of
40         integers
41         """
42         return self.img
43
44     def set_ans(self, ans):
45         """
46         Updates the answer for the image. Should be a single
47         integer between 1-4
48         """
49         self.ans = ans
50         return self
51
52     def get_ans(self):
53         """
54         Returns the answer as an integer between 1-4
55         """
56         return self.ans
57
58     def set_id(self, id):

```

```

57         """
58         Updates the id of the image
59         """
60         self.id = str(id).rstrip()
61         return self
62
63     def get_id(self):
64         """
65         Returns the id of the image
66         """
67         return self.id

```

../src/Image.py

```

1 import random
2 import math
3 from Utils import *
4
5 class Perceptron:
6     """
7     Perceptron represents a perceptron in the network. A
8     perceptron
9     represent one of the four feelings specified in the TYPE
10    field.
11    It reacts to an input image with a response based on it's
12    training.
13    Attributes:
14        TYPES: mapping of the perceptron type and type id
15        weights: a list of weights that maps to image pixels
16        type_id: the type id of the perceptron
17    """
18
19    def __init__(self, percept_type):
20        self.TYPES = {"HAPPY": 1, "SAD": 2, "MISCHIEVOUS": 3, "
21        MAD": 4}
22        self.weights = self._generate_weights(401)
23        self.type_id = self.TYPES[percept_type]
24
25    def process(self, img):
26        """

```

```

25         Uses the formula  $\sum(w[i], x[i])$  from lecture notes to
        process an image
        :return: summed weight of inputs
        """
27         l1 = apply(float.__mul__, img.get_img(), self.weights)
        return self._act(sum(l1))
29
31     def _act(self, val):
        """
        Computes the activation for the perceptron based on the
        sum of the
33         weighted input
        :param val: summed weight in input
35         :return: activation of the perceptron
        """
37         return 1 / (1 + math.exp(-val))
39
41     def update_weight(self, dw):
        """
        adds the provided delta error for the last iteration to
        each
        value in the internal weight matrix according to lecture
        notes
43         :return: None
        """
45         self.weights = apply(float.__add__, flatten(dw), self.
        weights)
47
49     def get_type_id(self):
        """
        Returns the type id for this perceptron. Consult the
        type
        type attribute for explanation of the meaning of the
        returned
51         value.
        :return: type id as an integer [1, 4]
53         """
        return self.type_id
55
57     def _generate_weights(self, acc):
        """
        Returns a list of the specified length containing pseudo
        random values in interval [0.0, 1.0)
59         :param acc: length of the list to generate
        :return: list of pseudo random floats
61

```

```

63         """
        return [random.random()] if acc == 1 else self.
_generate_weights(acc-1) + [random.random()]

```

../src/Perceptron.py

```

1 import random
from Utils import mmult, flatten
3
5 class Tutor:
    """
7     Tutor represents a supervisor object that is responsible for
    training a
    set of perceptrons using a provided set of training images
    until a
9     certain level of correctness is reached.
    Attributes:
11         perceptrons: a tuple of perceptrons to train
        images: a list of image objects to train on
13         learning_rate: value specifying how quickly the network
    should learn
        threshold: the value that indicates when to stop
    training
15     """
17     def __init__(self, perceptrons, images, learning_rate,
    threshold):
        self.perceptrons = perceptrons
19         self.images = images
        self.learning_rate = learning_rate
21         self.threshold = threshold
23     def train(self):
        """
25         Trains the perceptrons provided to this object with the
    provided
        using the provided training images until the desired
    threshold
27         for the sum squared error of the network is reached. The
    trained

```



```

29         perceptrons are returned as a tuple.
        """
        self._do_training([])
31     return self.perceptrons

33     def _do_training(self, errors):
        """
35         This function is what actually does what's described in
        the train
        method. This separation is done so that the user of this
        method is
37         not bothered by the fact that the method is implemented
        recursively.
        """
39         if self._accurate(errors): return self
        random.shuffle(self.images)
41         self._do_training(flatten(map(self._each_perceptron,
        self.images)))

43     def _each_perceptron(self, image):
        """
45         Shows the provided image to all the perceptrons and
        returns a list
        containing the error for every perceptron.
        """
47         return map(lambda p: self._expose(image, p), self.
        perceptrons)

49     def get_perceptrons(self):
        """
51         get_perceptrons returns a tuple of perceptrons. This
        method returns
53         the perceptrons in their current state. The train()
        method should be
        called until desired MSE is achieved before a call to
        this method
55         makes sense.
        :return: perceptrons as a tuple
        """
57         return self.perceptrons

59     def _expose(self, image, perceptron):
        """
61         Shows the provided image to the specific perceptron and
        uses the

```

```

63         formula for perceptron learning from lecture notes to
        update the
        internal weights in the perceptron based on the delta
        error
65         for the perceptron on the image.
        :return:
67         """
        err = self._error(self._desired(image, perceptron),
        perceptron.process(image))
69         dw = self._delta_weight(image, err)
        perceptron.update_weight(dw)
71         return err

73     def _delta_weight(self, inp, err):
        """
75         Uses the formula (error * learning rate * input) from
        lecture notes
        to produce the delta error for the activation of the
        perceptron
77         which is returned.
        :param inp: the perceptron input (image)
79         :param out: desired output, either 0 or 1
        :param act: the activation of the perceptron
        :return:
81         """
83         return mmult(inp.get_img(), self.learning_rate * err)

85     def _error(self, out, act):
        """
87         Computes the error between the desired output and the
        actual activation
        """
89         return out - act

91     def _accurate(self, err_list):
        """
93         Determines the size of the error using the formula for
        computing
        the error of a network provided over at:
95         http://lcn.epfl.ch/tutorial/english/perceptron/html/
        learning.html
        Uses the user provided threshold value to determine
        whether the
97         network performs well enough.
        """

```

```

99         if not err_list: return False
100         mse = (sum(map(lambda x: x**2, err_list)))/2
101         return mse < self.threshold
102
103     def _desired(self, img, percept):
104         """
105         Compares the provided image and perceptron to determine
106         whether
107         the neuron should be firing on the image. The decision
108         is
109         represented as either 0 or 1.
110         :return: 1 or 0 if the neuron should fire
111         """
112         return int(img.get_ans() == percept.get_type_id())

```

../src/Tutor.py

```

2
3
4     def mmult(m, c):
5         """
6         Receives a matrix and returns a new matrix
7         where every element in the matrix is multiplied
8         by the provided constant
9         :param m: matrix
10        :param c: constant
11        :return: new matrix with the applied constant
12        """
13        return map(lambda x: x * c, m)
14
15    def msum(m):
16        """
17        Returns the sum of every element in the matrix
18        """
19        return sum(map(sum, m))
20
21
22    def apply(op, l1, l2):
23        """
24        Returns a list of the same length as l1 and l2 where each

```

```

    index is the result of the provided operation on l1[i] & l2[
    i]
26     NOTE: Shadows built in function apply but since it is
    deprecated
    since version 2.3 we don't care.
28     Reference: https://docs.python.org/2/library/functions.html#
    apply
    """
30     return map(lambda tup: op(float(tup[0]), float(tup[1])), zip
    (l1, l2))

32
33 def flatten(l):
34     """
    Receives a nested list and returns all the values in a
    single list
36     Example: [[1], [2, 3], 4] -> [1, 2, 3, 4]
    :return: flattened list
    """
38     if not l: return l
40     return flatten(l[0]) + flatten(l[1:]) if isinstance(l[0],
    list) else l[:1] + flatten(l[1:])

```

../src/Utils.py